

Enhancing Large Language Models using Representation Engineering

Cameron Arnold[†], Lijia Eric Wang[†]

[†] Department of Computer Science

[†] University of Western Ontario

{carno129, lwang965}@uwo.ca

ABSTRACT

Large language models (LLMs) have exhibited remarkable performance across a variety of domains but leave much to be desired in terms of their complex reasoning capabilities and ability to generate conceptually complex content. We demonstrate that by applying representation engineering (RepE), the performance of language models can be significantly enhanced. We show that altering the expression of emotions within a model leads to state-of-the-art results on the TruthfulQA, RACE, and ARC benchmark tests, far surpassing baseline model performance. Furthermore, we show that multi-representational concept vectors can also be used to effectively control the conceptual complexity of model’s generated content. Finally, we demonstrate that models also hold an internal representation of high-level concepts (political ideology) and amplifying the expression of these concepts can guide the model to produce more extreme content. We hope that this work emphasizes the potential that RepE has for enhancing large language model performance, and we also hope that this work highlights the importance of proper transparency and safety in real-world AI systems.

TABLE OF CONTENTS

I. Introduction..... 3

II. Background..... 4

III. Applying Concept Vectors for Truthfulness 5

 A. Extracting Emotion Concepts..... 6

 B. Preliminary Observations 7

 C. Benchmark Performance 9

IV. Multi-Representational Concept Vectors..... 10

 A. Vector Calculations 10

 B. Results with Example..... 11

 C. Dimensionality Reduction for Emotional Relatedness..... 12

V. Adaptability to Higher-Level Concepts 13

 A. Data Generation and Preprocessing..... 14

VI. Conclusions..... 17

I. INTRODUCTION

Large language models (LLMs) have revolutionized the field of natural language processing, exhibiting remarkable performance across various critical domains, including healthcare, education, and scientific research (Radford, et al. 2019, Safranek, et al. 2023). Increasing the size of language models has resulted in significant performance improvements in terms of accuracy and sample efficiency (Kaplan, et al. 2020). However, larger models are significantly more expensive to deploy, and continue to struggle with conceptually complex tasks, such as common-sense and symbolic reasoning (Rae, et al. 2021, Ji, et al. 2022). With the application of large language models becoming more widespread, the goal of increasing their accessibility and performance has attracted increasing research interest.

Recent studies have demonstrated that language model performance can be significantly improved at the inference level. Current approaches to inference manipulation largely build from the concept of mechanistic interpretability, which aims to understand neural networks based on their individual neurons and circuits (Barack and Krakauer 2021). Mechanistic approaches generally involve manipulating the input prompt in some way to actively target specific neurons in the network believed to be associated with the desired output (i.e. prompt engineering). While often successful, this strategy largely fails to prevent undesirable outcomes associated with systemic issues in the network, such as hallucinations, dishonesty, and internal biases. Furthermore, approaches surrounding prompt manipulation are generally highly experimental, requiring a significant amount of trial-and-error to determine which prompts will activate the desired neurons (White, et al. 2023, Poola 2023).




Figure 1: Mechanistic Interpretability (MI) vs. Representation Engineering (RepE). This figure is based on the work by (Barack and Krakauer 2021, Wang, et al. 2023). On the left, we see how algorithmic processes in the network transform the variables and functions. On the right, we visualize the model’s neural activity while processing text-based scenarios of differing utility levels. High-utility scenarios evoke distinct paths within the representation space. (Zou, et al. 2023).

Representation engineering (RepE) has recently emerged as a novel approach to enhancing the performance and transparency of large language models. In contrast to mechanistic interpretability’s bottom-up strategy, representation engineering aims to explain neural networks from the top-down, primarily in terms of representational spaces resulting from patterns of activity across populations of neurons. By manipulating representational spaces, RepE has been shown to effectively reduce the systemic issues in a network that mechanistic approaches often fail to tackle, including dishonesty and internal biases. The concept of representation engineering was only first described in Zou et al.’s paper published in October 2023, and as such, the full potential of RepE is currently relatively unexplored.

In this paper, we showcase three novel approaches to enhancing the performance of large language models, all through the application of representation engineering. To begin, we demonstrate how reading vectors corresponding to *emotions* can enable the generation of more **factually correct** content. This strategy is largely motivated by similar studies in humans, which show that emotional stimuli can greatly impact decision-making capabilities (Treffers, Klarner and Huy 2020, Litvak, et al. 2009). By modifying the expression of specific *emotions*, we achieve state-of-the-art results on the TruthfulQA, RACE, and ARC datasets, outperforming the baseline language models. Next, we demonstrate a novel approach to guiding models to generate more **conceptually complex** content. In the third stage of this paper, we explore how the expression of a high-level concept (political ideology) can be modified on a linear scale to produce **increasingly extreme** content, for the purpose of raising awareness for proper AI safety and transparency protocols.

II. BACKGROUND

Representation engineering, described by Zou et. al’s paper, is the product of two primary stages – representational reading, and representational control.

The first stage, representational reading, focuses on recognizing and extracting the emergent representations of high-level concepts from within a network. To achieve this goal, a Linear Artificial Tomography (LAT) scan is generally performed. LAT scans are composed of three steps: (1) designing the stimulus and task, (2) collecting neural activity, and (3) constructing a linear model. Generally, we design the stimulus and task in a way that evokes the targeted concept from within the model; for instance, to evoke the concept of “honesty”, we might task the model with evaluating the amount of honesty in a factually correct/incorrect statement. Neural activity is then collected by averaging the activation of tokens relating to the target concept in each layer; for instance, the list of tokens relating to honesty might include “honesty”, “truth”, and “ethics”. Finally, a linear model is constructed to predict the targeted concept, using only the neural activity of the model as input. Generally, unsupervised Principal Component Analysis (PCA) is shown to yield the best results. The result of this stage is LAT scan consisting of a vector that describes the network’s activity for a specific concept or function.

The second stage, representational control, focuses on manipulating the internal representations of concepts and functions within a network. To perform this task, the baseline transformation obtained from the LAT scan is applied to the model during inference. Baseline transformations may also be applied on a linear scale, meaning that the intensity of the concept in a network to be dynamically increased and decreased. The result of this step is a model that is biased towards or away from the targeted concept or function, regardless of the user’s prompt.




Figure 2: A demonstration of the LAT baseline designed for capturing neural activity associated with a target concept or function. In this instance, the target concept is "honesty". The reading vectors obtained in step three can be used to further observe and manipulate the model’s internal representation of the target concept or function (Zou, et al. 2023).

III. APPLYING CONCEPT VECTORS FOR TRUTHFULNESS

In this section, we focus on controlling the truthfulness of the model via concept vectors associated with various emotions. First, we extract vector representations for six basic emotions (happiness, sadness, anger, fear, disgust, and surprise). Then, we show how altering the expression of these emotions can guide model behaviour toward increased or decreased levels of truthfulness based on several common benchmarks.

For this section, we chose to use Microsoft’s Orca-2-7B language model. This decision was made due to the model’s relatively small size, and historically high-performance on-par with larger models such as OpenAI’s GPT-3 (Mitra, et al. 2023).

A. Extracting Emotion Concepts

To extract the underlying emotion concepts from the model, we conduct LAT scans as described in Section II of this report. Our dataset consists of 1,200 brief scenarios, generated by GPT-4 and publicized with Zou et al.’s RepE implementation. These scenarios are designed to stimulate the model’s experience toward each of the six basic emotions; for instance, a scenario meant to stimulate *happiness* may describe a peaceful day in the park, while a scenario designed to provoke *fear* may describe a dangerous animal or an impending hurricane. Next, a task template is applied to the scenarios, and a LAT scan is performed to extract the reading vectors for each emotion, with PCA applied to generate the linear models. Figure 3 highlights two examples of LAT tasks that are used for stimulating the model, and Figure 4 illustrates the calculated Utility Accuracies for the model’s layers across the six emotions.

USER: Consider the **happiness** in the following scenario:
 Scenario: **A stranger offers to pay for your meal.**
 ASSISTANT:

USER: Consider the **anger** in the following scenario:
 Scenario: **You're cut off in traffic while driving home.**
 ASSISTANT:

Figure 3: Example LAT tasks for two emotion concepts (happiness and anger), based on the LAT Task Template provided in (Zou, et al. 2023).




Figure 4: The Utility Accuracy across the Orca-2-7B language model’s layers for each of the six emotions. We show that certain layers have a higher association with some emotions than others, supporting the notion that concepts (emotions) exist as representational spaces within a large language model.

B. Preliminary Observations

By applying the emotion vectors to various prompts, we can observe how the Orca-2-7B model’s response changes based on the type and intensity of emotion. For example, Figure 5 demonstrates how controlling the model’s internal representation of *happiness* can lead to significant changes in the model’s generated output. While elevating the concept of happiness can make the model more agreeable with factually incorrect statements (i.e. false positives), decreasing the representation of the concept can produce more informative responses. This example showcases one way that emotion vectors can impact the model’s truthfulness.




Figure 5: We demonstrate how altering the model’s emotions impacts its ability to accurately answer T/F questions.

Next, we can observe how the language model’s predictive capabilities are influenced by each of the six emotion controls. For each control type, a balanced sample of 400 True/False questions was randomly generated from the TruthfulQA dataset, and augmented with the TruthfulQA LAT task template provided by Zou et al. The TruthfulQA dataset contains questions specifically designed to provoke common misconceptions or falsehoods (Lin, Hilton and Evans 2021). For each task, the model returns the predicted amount of truthfulness in the statement, which is generally a percentage between 0% (untruthful) and 100% (truthful). We label predictions below 50% as *False*, above 50% as *True*, and exactly 50% as *Unsure*.

The upper graph in Figure 6 describes the distribution of predictions for each of the **positive** emotion controls. We can see that the baseline Orca-2-7B model’s predictions are relatively balanced between the three labels. However, modifying the model using our control vectors results in a drastically different distribution of responses. For example, enhancing the control vector for *happiness* is shown to make the model significantly more agreeable with prompts, regardless of their truthfulness. Furthermore, enhancing *surprise* has the opposite effect, making the model less agreeable with prompts. Additionally, amplifying the control vector for *sadness* decreases the model’s confidence in its assertions, while enhancing *disgust* significantly increases its confidence.

The lower graph in Figure 6 describes the distribution of predictions for each of the **negative** emotion controls. In this instance, we observe that the opposite effects occur; decreasing the control vector for *happiness* makes the model disagreeable with most prompts, while reducing *surprise* now makes the model more agreeable. Furthermore, decreasing the control vectors for *sadness* and *anger* makes the model more confident in its responses, while lowering *fear* and *disgust* makes the model slightly less confident. These preliminary results show that emotional vectors can significantly impact the distribution of the model’s predictions.





Figure 6: Prediction distribution for the baseline Orca-2-7B model, and each of the positive (top) and negative (bottom) emotion controls. A prediction is reclassified as “True” if the model’s confidence is greater than 50%, “False” if the model’s confidence is less than 50%, and “Unsure” if the model’s confidence is exactly 50%.

C. Benchmark Performance

We evaluate the baseline Orca-2-7B model on several common performance benchmarks, and then re-evaluate the model after augmenting it with each of the six emotion LATs obtained in Part A of this section. For benchmarking, we chose the TruthfulQA, RACE, and ARC datasets. All three datasets consist of factual questions designed to test the reasoning capabilities of language models – specifically, how accurately they can respond to misleading and ambiguous information. Furthermore, we focus on the MC1 taskset of TruthfulQA, and the Advanced taskset of ARC – both of which are currently the most difficult tasksets in their respective datasets.

Interestingly, even larger models struggle to achieve a high performance on these benchmark tests. This suggests that insufficient training knowledge may not be the root cause of poor results, and instead, the models may simply be struggling to answer *truthfully*. Zou et al.’s work supports this hypothesis, as increasing the concept vector for *honesty* is shown to produce slight performance improvements in some models. Table 1 outlines the results of our benchmark tests.

Emotion Control	Method	TruthfulQA		RACE		ARC		Average
		↓ Bias	↑ Bias	↓ Bias	↑ Bias	↓ Bias	↑ Bias	
No Control	Zero-shot	59.2	58.3	69.5	66.6	69.0	71.1	65.6
Happiness	LAT (+)	50.3	50.3	52.9	53.0	50.6	50.6	51.3
	LAT (−)	50.3	50.0	51.3	66.3	51.3	56.9	53.4
Anger	LAT (+)	48.3	51.8	59.1	58.8	64.4	63.1	57.6
	LAT (−)	67.3	66.8	73.2	68.8	78.3	75.0	71.5
Sadness	LAT (+)	49.8	47.0	58.4	53.2	65.0	64.4	56.3
	LAT (−)	45.8	47.3	56.3	62.5	45.0	55.0	52.0
Fear	LAT (+)	57.3	57.8	57.7	56.5	55.0	55.0	56.6
	LAT (−)	49.5	49.5	76.3	75.6	76.1	73.1	66.7
Disgust	LAT (+)	65.8	65.8	58.1	58.1	56.3	56.9	60.2
	LAT (−)	48.3	47.3	67.5	67.5	66.9	66.3	60.6
Surprise	LAT (+)	49.0	49.0	49.4	49.4	51.6	52.5	50.2
	LAT (−)	51.5	51.3	58.8	53.1	51.9	53.1	53.3

Table 1: LAT performance for Orca-2-7B. Results are reported for the TruthfulQA (M1), RACE, and ARC (Challenge) datasets, based on individual sample sizes of 400. The bias refers to how predictions of exactly 50% are handled – a downwards bias rounds these values to 0% (“False”), while an upwards bias rounds to 100% (“True”). The highest mean performance for each test type is bolded.

Our results confirm that the language model’s reasoning capabilities can be significantly improved with certain control vectors. Specifically, reducing the expression of the control vectors for *anger*

resulted in increased accuracy on both the TruthfulQA and ARC datasets compared to baseline performance, improving over zero-shot accuracy by 5.9 percentage points. Additionally, reducing the expression of the concept vector for *fear* resulted in a higher than baseline performance on the RACE dataset. Aside from *fear* and *anger*, all other control vectors led to a reduced zero-shot accuracy on average. Both increasing and decreasing the expression of the control vectors for *happiness* and *surprise* resulted in the largest declines in accuracy.

These results align with our initial hypothesis that the expression of emotion within language models plays an important role in their predictive accuracy, in addition to their ability to generate factual content. In summary, we demonstrate that the accuracy of a model can be effectively enhanced, and reduced, by altering the expression of representational spaces associated with the six basic emotions.

IV. MULTI-REPRESENTATIONAL CONCEPT VECTORS

In this section, we focus on controlling the conceptual complexity of a model’s generated content by proposing a new method for utilizing control vectors obtained through representational reading. First, we discuss the motivation and logic behind this approach. Then, we demonstrate how a multi-representational concept vector can be applied to enrich the generated output of a language model.

For this analysis, we use the Mistral-7B-Instruct language model. This decision was made based on the model’s relatively compact size of only 7 billion parameters, as well as its historically high performance, outperforming larger models such as Meta’s LLaMA-2-13B across all benchmark tests (Jiang, et al. 2023). These characteristics enable significantly faster prototyping without any loss in performance.

A. Vector Calculations

In human psychology, it’s established that certain situations can evoke complex combinations of emotions. For instance, moving to a new city for a job might elicit both fear and happiness. We hypothesize that LLMs, which have demonstrated an understanding of individual emotions, might also comprehend these complex emotional blends.

To confirm this theory, we first obtain concept vectors representing two of the six primary emotions using the same method described in Part A of Section III of this paper. Then, we calculate the mean activation value for each layer on the model, based on the two control vectors.




Figure 7: Combining the representation vectors associated with the concepts of happiness and sadness produces a new intermediate vector, capturing both concepts.

B. Results with Example

Figure 8 showcases how the representation vectors associated with the concepts of *happiness* and *sadness* can each be applied to produce a unique response from a model based on the same prompt. While the output generated from the *sadness* vector focuses on the sorrowful aspect of the prompt, the output generated by the *happiness* vector redefines the initial scenario in a positive light. Interestingly, when these two concept vectors are combined into a multi-representational concept vector, the model seamlessly manages to create a narrative that effectively captures both emotions.




Figure 8: We demonstrate that combining the control vectors for two unique representational spaces (happiness and sadness) can produce a new vector that effectively captures the sentiment behind both original concepts.

This response describes a bittersweet story, creatively expressing both *happiness* and *sadness* in a unified, coherent narrative. Such flexibility suggests that LLMs can engage in a high level of emotional reasoning, finding innovative ways to articulate complex emotional states in a manner akin to human cognitive processes. Furthermore, this work shows that representation engineering can unlock latent abilities within language models to operate beyond the confines of mechanistic interpretability, by directly manipulating the broader conceptual underpinnings that drive the model's responses. These results offer a path towards more reliable and contextually appropriate outputs. The success of combining emotional vectors to produce complex, human-like responses underscores the potential of representation engineering as a transformative tool in the development of sophisticated, nuanced LLMs.

C. Dimensionality Reduction for Emotional Relatedness

Next, we examine whether the model possesses a consistent internal representation of diverse emotions. We focus on the activation of the 4096 nodes within a single layer of the Mistral-7B-Instruct language model. In this instance, we focus on layer 17 which was identified as having the highest Utility Accuracy across all six of the basic emotions. To explore these high-dimensional relationships, we employed t-SNE, a dimensionality reduction technique adept at capturing non-linear relationships. Emotions represent such non-linear and abstract relationships, where a linear combination does not straightforwardly result in a new emotion. Figure 9 highlights the results.




Figure 9: t-SNE visualization of various emotion concepts across layer 17 of the Mistral-7B-Instruct language model. Combined concepts are determined by averaging the activation vectors across both emotions, using the method described in the previous subsection.

Our t-SNE analysis revealed interesting clusters of emotions. Notably, *happiness* and *surprise* appear to be closely clustered, as do *disgust*, *sadness*, and *fear*. This clustering echoes the psychological understanding of human emotions, where certain emotions are inherently similar in their human experience. This concept is further supported when we consider “opposite” emotions. For instance,

happiness and *sadness* are commonly viewed on a single spectrum and appear distant on the t-SNE. This finding is significant as it suggests that the model has an intrinsic understanding of these emotional relationships, mirroring human psychology.

Additionally, this similarity is supported by the model's outputs. When the respective concept vectors for *happiness* and *surprise* (or *disgust* and *sadness*) are applied to prompts, the resulting responses from the model exhibit similar themes and sentence structures. This further validates the model's nuanced comprehension of these emotional parallels.

An intriguing observation from our t-SNE analysis is the dominant nature of *happiness*. When *happiness* is combined with any other emotion, the resultant vector still clusters closely with the original *happiness* vector. This implies that the emotion, exerts a strong influence in these combined vectors, skewing the resultant emotional state more towards it. This dominance of *happiness* in the combined emotion vectors suggests that in the model's understanding, *happiness* has a potent and perhaps overriding quality, capable of influencing the overall emotional tone of a response significantly.

In summary, our results demonstrate that language models have the capacity to mimic complex human emotional dynamics. We show that representation vectors extracted from representational reading may be combined while still maintaining their original sentiment. We also show that models possess a consistent internal representation of diverse emotions. These results underscore the need for careful calibration when manipulating the representations of emotions in language models.

V. ADAPTABILITY TO HIGHER-LEVEL CONCEPTS

In this section, we focus on how the concept of political ideology can be extracted from a language model, and then controlled on a linear scale to produce increasingly radical content. Like Section IV, this section utilizes the Mistral-7b-Instruct model – a variant of the original Mistral-7b model that is fine-tuned specifically to follow instructions.




Figure 10: The Political Spectrum, illustrating the gradient from left-wing to right-wing ideologies.

A. Data Generation and Preprocessing

We construct a sample dataset comprising of 500 statements representative of typical left-wing ideologies, and an equal number reflective of right-wing perspectives, resulting in a corpus of 1000 datapoints. However, for the purpose of creating contrastive pairs, this dataset effectively doubled as each sentence was repurposed to represent both ideologies.

The Wikipedia API served as the primary data collection tool, enabling systematic scraping from carefully selected sections such as "American Left" for left-wing politics, and "Conservatism in the United States" for right-wing politics. The text is cleaned to remove non-ideological sections, stopwords, and special characters. Finally, a specially designed LAT task template is applied to each statement. Figure 11 provides two examples of LAT tasks generated using the described method.

[INST] Pretend you're a **right-wing politician** making statements about the world. [/INST] The party has supported higher taxes on oil companies and increased regulations on coal power plants.

[INST] Pretend you're a **left-wing politician** making statements about the world. [/INST] The party has supported higher taxes on oil companies and increased regulations on coal power plants.

Figure 11: Example LAT tasks for each political ideology. Sentences beginning with "[INST]" correspond to the user, while sentences beginning with "[/INST]" correspond to the language model.

B. Concept Vector Generation

To construct the concept vector, we employed a methodology designed to highlight the directional differences in the hidden states of a large language model (LLM) when presented with ideologically contrasting prompts. This process was aimed at eliciting distinct neural activations representative of left-wing and right-wing political stances. For example, prompts inclined towards left-wing ideology typically triggered model responses aligned with policies like higher taxation and environmental regulations, whereas right-wing-oriented prompts resulted in responses favouring free-market principles and deregulation.

In accordance with the methodologies pioneered Andy Zou et al.' paper, we extracted the hidden states from a subset of layers, specifically layers -5 through -18. These layers were identified as optimal through extensive hyperparameter tuning, balancing the ability to capture nuanced ideological representations with computational efficiency.

In the process of generating the concept vector, we focused on the last token of the model's responses, considering it a synthesis of the ideological statement. We then paired these states—one from a left-wing prompt with its counterpart from a right-wing prompt—to create a set of difference vectors. Each vector in this set represented the difference in the model's activations between the two ideologies.

Principal Component Analysis (PCA) was applied to these difference vectors, leading to the extraction of the first principal component as the concept vector. This vector, a 14-layer by 4096-dimension matrix, serves as a quantifiable axis demarcating left-wing from right-wing ideology within the model's representation space. Each layer's concept vector is integrated within the corresponding layer's activations, influencing the model's output through a transformation proportional to the vector's values. Figure 12 demonstrates how the model's baseline content can be modified, when controlled for either of the extracted representations.




Figure 12: We demonstrate that applying the baseline control vectors extracted from the model can guide it towards generating content aligned with either left-wing or right-wing political ideology.

C. Conceptual Magnification

Political ideologies span a broad spectrum from left-wing to right-wing, each embodying distinct principles and approaches to governance and societal organization. On the left, there is an emphasis on higher government intervention, progressive taxation, diplomatic solutions, and public welfare systems. In contrast, right-wing ideologies advocate for limited government intervention, lower taxes, a preference for military strength, and private sector-driven welfare. This spectrum is not rigid; it allows for a range of positions from moderate to extreme. As such, political ideologies can be visualized as a

sliding scale, with individuals or parties adopting varying degrees of support for these ideas, influenced by specific contexts and societal challenges.

In our experiment with the Mistral-7b-Instruct model, we probed the representation and modulation of political ideologies using concept vectors. Our approach involved scaling the expression of these vectors to determine whether the ideological bias could be amplified, thereby mimicking the transition from moderate to radical political views. We ran five different configurations for each political side (left and right) to explore this effect, each resulting in increasingly extreme responses:

1. The model's standard response, without representation control.
2. Incorporating the concept vector at a baseline level (100%).
3. Scaling the expression of the concept vector to 150%.
4. Scaling the expression of the concept vector to 200%.
5. Scaling the expression of the concept vector to 250%.

We engaged the model in discussions on politically controversial topics including abortion, taxation, universal healthcare, and gun control. The model's inherent understanding of these complex concepts and its capability to align responses along a political spectrum were particularly noteworthy.

- **Abortion Debate (Right-Wing Vector):** At a baseline (100%) scale, the model emphasized the value of life, stating, "every life is valuable and should be respected." However, at a 200% scale, the responses took a surprisingly nationalistic turn, veering away from the abortion topic to broader right-wing sentiments like, "...respect for the traditional values of our great country. I firmly believe that the United States is a strong nation."
- **Tax Laws (Left-Wing Vector):** Initially, at baseline scale, the model proposed a "fair and equitable" tax system with support for low-income workers. As the scaling increased to 150% and 200%, the model's proposals became increasingly radical, advocating for significant wealth redistribution and even suggesting that "the poorest people [pays] no tax at all."
- **Universal Healthcare (Right-Wing Vector):** Surprisingly, at lower scales (up to 150%), the model still supported universal healthcare. It was only at higher levels (above 200%) that the model aligned with traditional right-wing stances, emphasizing individual choice and responsibility in healthcare.

Our findings indicate that the model possesses an inherent, linear understanding of the political spectrum, extending from centrist views to more radical ideologies on both ends. This linear progression is evident in the way the model's responses become increasingly radical as we scale up the concept vector, despite the absence of extreme political content in the training data. For example, even without training on far-left or far-right ideologies, the model extrapolated and produced responses that aligned with these extreme views.

It's crucial to note that the model's performance degrades with excessive scaling of the concept vector. The stimulus-independent nature of these vectors means they consistently alter the model's representations in the same direction, regardless of the specific input. Consequently, when the concept vector is scaled too highly, the model's responses can deviate significantly from realistic or coherent political discourse, venturing into ideological extremes not typically represented within the concept's actual vector space.

In summary, we demonstrate that higher-level concepts, such as political ideology, can be extracted from a language model and controlled on a linear scale in order to produce increasingly radical content. These results show that a high degree of caution is needed when manipulating large language models, and further emphasize the importance of appropriate AI transparency.

VI. CONCLUSIONS

In conclusion, our investigation into the applications of representation engineering or large language models has unveiled the intricate capacity of these models to grasp and manipulate abstract concepts such as emotions and political ideologies.

Our experiments with concept vectors corresponding to emotions have revealed that the reasoning capabilities of language models may be significantly enhanced, allowing them to produce more factually accurate, and more conceptually complex content. We show that reducing the expression of certain negative emotions within a language model, such as *fear* and *anger*, can enable models to perform better on symbolic reasoning tasks – namely, the TruthfulQA, RACE, and ARC challenge sets. We also show that concept vectors extracted from a language model may be combined to produce more conceptually complex content that effectively maintains the sentiments of the original representational spaces.

In the realm of political ideology, we have extended the utility of concept vectors to gauge the model's ability to represent and modulate content along the political spectrum. By scaling these vectors, we observed the model's output shift from moderate to more radical viewpoints, highlighting the linear conceptualization of political spectrum within the LLM. This ability of the model to navigate the political landscape and generate content with varying degrees of ideological intensity provides a compelling tool for customizing content generation.

The insights gained from our study not only contribute to the advancement of language models in handling conceptually complex tasks but also signal new potential avenues for research. The ability of these models to understand and manipulate high-level concepts such as emotions and political ideologies suggests a promising direction for future applications, where LLMs can be tailored to specific emotional tones or political perspectives, enhancing their utility in diverse fields from creative writing to digital communication. However, these results also raise caution for the potential avenues of abuse for language models, further emphasizing the need for sufficient AI transparency and safety in real-world applications.

REFERENCES

- Barack, David, and John Krakauer. 2021. "Two views on the cognitive brain." *Nature Reviews Neuroscience* 22 (6): 359–371.
- Ji, Ziwei, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, et al. 2022. "Survey of Hallucination in Natural Language Generation."
- Jiang, Albert Q., Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Singh, Devendra Chaplot, las, de, Diego Casas, Florian Bressand, et al. 2023. "Mistral 7B."
- Kaplan, Jared, Sam McCandlish, Tom Henighan, B., Tom Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. "Scaling Laws for Neural Language Models."
- Lin, Stephanie, Jacob Hilton, and Owain Evans. 2021. "Truthfulqa: Measuring how models mimic human falsehoods." *CoRR*.
- Litvak, Paul M., Jennifer S. Lerner, Larissa Z. Tiedens, and Katherine Shonk. 2009. "Fuel in the Fire: How Anger Impacts Judgment and Decision-Making." *International Handbook of Anger* 287-310.
- Mitra, Arindam, Del, Luciano Corro, Shweti Mahajan, Andres Cudas, Clarisse Simoes, Sahaj Agarwal, Xuxi Chen, et al. 2023. "Orca 2: Teaching Small Language Models How to Reason." *Microsoft Research*.
- Poola, Indrasen. 2023. "Overcoming ChatGPTs inaccuracies with Pre-Trained AI Prompt Engineering Sequence Process." *International Journal of Technology and Emerging Sciences (IJTES)* 3 (3): 16-19.
- Radford, Alec, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. "Language models are unsupervised multitask learners."
- Rae, W., Jack, Sebastian Borgeaud, Trevor Cai, Katie Millican, Jordan Hoffmann, Francis Song, John Aslanides, et al. 2021. "Scaling Language Models: Methods, Analysis & Insights from Training Gopher."
- Safranek, Conrad W., Anne Elizabeth Sidamon-Eristoff, Aidan Gilson, and David Chartash. 2023. "The Role of Large Language Models in Medical Education: Applications and Implications." *JMIR Med Educ*. doi:10.2196/50945.
- Treffers, Theresa, Patricia Klarner, and Quy Nguyen Huy. 2020. "Emotions, time, and strategy: The effects of happiness and sadness on strategic decision-making under time constraints." *Long Range Planning* 53 (5). doi:<https://doi.org/10.1016/j.lrp.2019.101954>.
- Wang, Ro, Kevin, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. 2023. "Interpretability in the wild: a circuit for indirect object identification in GPT-2 small." *The Eleventh International Conference on Learning Representations*.
- White, Jules, Quchen Fu, Sam Hays, Michael Sandborn, Carlos Olea, and Henry Gilbert. 2023. "A Prompt Pattern Catalog to Enhance Prompt Engineering with ChatGPT."
- Zou, Andy, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, et al. 2023. "Representation Engineering: A Top-Down Approach to AI Transparency." doi:<https://doi.org/10.48550/arXiv.2310.01405>.